

NLP Meets the World: Companion Notes

Shomir Wilson

Remarks

I'm sharing some behind-the-scenes notes about this paper to explain why I wrote it, how it came together, and how reviewers' comments improved it. My motivation for these notes is to benefit other researchers, especially students who are new to writing papers for publication. I'm also motivated by an enjoyment of explaining the hidden systems of academia and research. You can read more of my academic guidance and experience reports at <https://shomir.net/advice>.

As a Ph.D. student, my first paper submission also was a reflection paper. I submitted it to a small conference—Commonsense 2009—and it received one very enthusiastic review and two very negative ones; it was rejected. After that I avoided writing reflection papers for a long time, and getting this paper accepted to EMNLP is, in a way, the conclusion of a 17-year story arc. It's also the most unusual paper I've submitted to an NLP conference: it contains no algorithms, experiments, data, or math. It's readable by a general audience.

This enhanced version of the paper lives at <https://shomir.net/nlpmw>, where you can also find a link to the published version on ACL Anthology. Remember the published version is appropriate for citing and referencing, not this one, unless you're specifically referring to these companion notes.

Below is the commentary that pairs with the highlighted text. I think research writing can have depth and nuance, and although it's a paper with a sincere message, these callout notes also reveal some metatextual humor. After that is a timeline that shows the sequence of events that led to this paper's acceptance by EMNLP.

Commentary

1 I wrote this precolon as an allusion to a 1990s TV show called "Boy Meets World", which I never watched. Conveniently, readers in 2026 might know the similarly-titled Netflix series "Earle Meets World", which I've also never watched.

2 The observations in this paper are broadly applicable to AI, but I targeted NLP to appeal to EMNLP reviewers.

3 "For individual researchers" is an important nuance here: the interests of the NLP research community as a whole are connected to the interests of researchers within it. Some competition is inevitable in research, but if you're talking with news media then you're probably experiencing success. It's a time when supporting other people is easy and benefits both you and them.

4 The first submission of this manuscript instead addressed improving conversations "with the public", but a reviewer criticized that as too vague.

5 A reviewer suggested adding a visual distinction between citations to research and citations to news articles. In practice that distinction leaves out some sources, and instead I marked sources that weren't peer reviewed (to the best of my knowledge). Note some citations to non-peer reviewed sources still command respect from the research community, such as the two research communication frameworks this manuscript builds upon.

6 I avoid evaluating this claim (and many others) in this manuscript, but it's among the boldest claims I've ever paraphrased or quoted in a published paper.

7 Three of my favorite under-appreciated older movies that portray anxieties about AI: (1) *Colossus: The Forbin Project* (1970), in which a super-computer enacts an agenda of totalitarian world rule through control of nuclear weapons; (2) *Electric Dreams* (1984), in which young man's personal computer becomes sentient and a romantic rivalry ensues; and (3) *A. I. Artificial Intelligence* (2001), in which a robotic child goes on a quest to become human to earn a mother's love.

8 All AI and NLP researchers are diplomats for their research communities, whether they're high schoolers talking with their parents or distinguished professors on live TV.

9 "Dear reader, I know this is a strange paper, but could you cite it?"

10 Connecting to these frameworks was the biggest difference between this manuscript and the first submission. Reviewers wanted evidence that the particular challenges I chose were meaningful.

11 This is the *section sign* (look it up in Wikipedia), and it's an underused space-saver in NLP papers.

12 I'd be glad to see (and even speak at) a science communication tutorial at an ACL event.

13 "Dear reviewer, conferences have published strange papers like this before. It's OK!"

14 A reviewer of the first version of this manuscript misunderstood §3.1 as supporting or refuting cognitive claims about LLMs, which it doesn't. I added a footnote to make that lack of stance explicit. This section broadly reflects my interests in this area, which are less about what constitutes LLMs "thinking" or "understanding" and more about the consistency or inconsistency of claims using these terms.

15 In other words: imagine two networks of entailments between word meanings, one representing NLP researchers' notions and the other representing the general public's notions. It's easy to assume the networks are identical, but they aren't. This is a graph isomorphism argument, though I didn't realize it when I wrote it.

16 This is one of a couple self-citations that appear in the paper. Junior researchers sometimes doubt whether it's OK to self-cite in a manuscript they will submit for double-blind review. It's actually fine if you remember to refer to your prior publications in the third person, as you would anyone else's work.

17 I originally wrote this observation in a social media post, and then I wanted to write an entire paper about it, but it wasn't enough. The idea of writing a paper with advice on research communication came shortly after.

18 Incredibly, this felt necessary to say. It doubles as dry wit.

19 I've occasionally heard NLP researchers argue against using cognitive terms for LLM behavior and then, a few minutes later, use them for that. I think we (researchers or anyone discussing AI) need to acknowledge the naturalness of metaphor while being fully aware of its limitations.

20 Be supportive and sympathetic even ("I know, it's easy to slip into talking about LLMs as *thinking*"), but make sure people understand the limitations of the language we use.

21 "[R]adical improvements" and "global peace": these claims reflect the unprecedented public atten-

tion to AI and NLP. Note that the cited Amodei essay received lots of press coverage, making it appropriate to mention.

22 Very rarely are students in my classes familiar with the AI winters before I introduce them. It's history we should learn from rather than repeating.

23 Apropos remembering AI history, this quote from the Lighthill Report remains remarkably prescient of contemporary critiques of AI.

24 There's an old saying that *AI* is a term for computing-related things that don't work well yet. This is a corollary to that claim.

25 NLP research has accomplished some great things. Remembering that is good for research community morale, not just research communication.

26 If you've never used Eliza, which is widely recognized as the first chatbot, try an online demo of it. After using it, consider how easy it is to anthropomorphize a chatbot that's nearly equivalent to a hundred-sided die with sentences on its sides.

27 This sentence was the direct result of a reviewer comment. Researchers' ethics aren't necessarily absent in situations like this; sometimes they're just different (and the researchers are unaware of the difference).

28 The failures in this paragraph came to mind more easily than successes. While the failures deserve attention, I sometimes think the NLP research community forgets how much it's accomplished.

29 This reflects what I've noticed in casual conversations about AI with people outside of computing research or practice. They often have a vague, worried question that resembles "Do you think AI is good?" or "Is AI harming humanity?" Questions like those reveal the opacity of AI to a general audience, as if AI is one thing instead of many.

30 Human-centered NLP research deserves more promotion to the public, for its impact and also for its value toward engaging public goodwill.

31 "Dear reviewer, accept this paper!"

32 Many PhD programs require a "how to do research" class, and I think research communication to the public is a topic that belongs in it. The relevance goes beyond talking with reporters; junior researchers often struggle to communicate their work even to other researchers in the same area.

33 There are limits to what a researcher can do to promote their work. Thanks to these profession-

als for the opportunities they’ve created for me to speak with news media about AI and NLP.

34 NSF awardees, remember to do this.

35 “Dear reviewer, this paper existing is better than this paper not existing, right?”

Timeline

Briefly, I submitted the manuscript twice to ACL Rolling Review and committed it three times to conferences. The first two commitments were declined and the third was accepted.

June 8, 2026: I posted on social media some thoughts about using analytic philosophy to evaluate the language used in cognitive claims about LLMs’ abilities. Soon afterward I began thinking about expanding the ideas into a publishable paper.

June 15, 2025: I began writing the paper, but I realized that I didn’t have enough to say about the initial topic. I decided to write more broadly about the challenges of communicating NLP research to the general public.

July 14, 2025: I submitted the manuscript to ACL Rolling Review well ahead of the July 28 deadline. I had overestimated the time I needed to write.

September 2025: The reviews released for the author response period. The paper received overall scores of 2, 4, and 1.5. (ACL Rolling Review’s scale for overall scores is 1–5, with 5 being the strongest.) The reviews contained some helpful suggestions, but the one with the lowest score contained a provocative comment about race and gender, which I reported to the metareviewer. The metareviewer assigned the paper an overall assessment of 4.

October 3, 2025: I committed the paper plus reviews to IJCNLP-AAACL 2025.

October 25, 2025: IJCNLP-AAACL declined it.

March 7, 2026: I committed the paper plus reviews to ACL 2026.

April 7, 2026: ACL declined it.

May 17, 2026: Having made substantial changes to the paper, I submitted a revision to ACL Rolling Review. I requested different reviewers but the same metareviewer.

July 2026: The reviews released for the author response period, and the paper received overall scores of 4, 2.5, and 4. The 2.5 review spoke positively of the paper and listed some “minor edits”. I replied to that review, agreeing with most of the edits point-by-point, and remarked that the overall score seemed low compared to their criticisms. The

reviewer raised it to 4. The metareviewer gave the paper an overall score of 4.5.

August 1, 2026: I committed the paper plus reviews to EMNLP 2026.

August 21, 2026: EMNLP accepted the paper.